



JAKARTA EE

Jakarta Data

1.0.0-b1, December 29, 2022: Draft

Table of Contents

Copyright	2
Eclipse Foundation Specification License	2
Disclaimers	2
Jakarta Data	4
1. Introduction	5
1.1. Goals	5
1.2. Non-Goals	5
1.3. Conventions	5
1.4. Jakarta Data Project Team	5
1.4.1. Project Leads	5
1.4.2. Committers	6
1.4.3. Mentor	6
1.4.4. Contributors	6
2. Repository	7
2.1. Repositories on Jakarta Data	7
2.2. Entity Classes	8
2.2.1. Entity Property Names	9
2.3. Queries Methods	9
2.3.1. Using the Query Annotation	9
2.3.2. Query by Method	10
2.4. Special Parameter Handling	12
2.5. Keyset Pagination	13
2.5.1. Example of Appending to Queries for Keyset Pagination	14
2.5.2. Avoiding Missed and Duplicate Results	14
2.5.3. Restrictions on use of Keyset Pagination	15
2.5.4. Keyset Pagination Example with Sorts	15

Specification: Jakarta Data

Version: 1.0.0-b1

Status: Draft

Release: December 29, 2022

Copyright

Copyright (c) {inceptionYear} , {currentYear} Eclipse Foundation.

Eclipse Foundation Specification License

By using and/or copying this document, or the Eclipse Foundation document from which this statement is linked, you (the licensee) agree that you have read, understood, and will comply with the following terms and conditions:

Permission to copy, and distribute the contents of this document, or the Eclipse Foundation document from which this statement is linked, in any medium for any purpose and without fee or royalty is hereby granted, provided that you include the following on ALL copies of the document, or portions thereof, that you use:

- link or URL to the original Eclipse Foundation document.
- All existing copyright notices, or if one does not exist, a notice (hypertext is preferred, but a textual representation is permitted) of the form: "Copyright (c) [\$date-of-document] Eclipse Foundation, Inc. <<url to this license>>"

Inclusion of the full text of this NOTICE must be provided. We request that authorship attribution be provided in any software, documents, or other items or products that you create pursuant to the implementation of the contents of this document, or any portion thereof.

No right to create modifications or derivatives of Eclipse Foundation documents is granted pursuant to this license, except anyone may prepare and distribute derivative works and portions of this document in software that implements the specification, in supporting materials accompanying such software, and in documentation of such software, PROVIDED that all such works include the notice below. HOWEVER, the publication of derivative works of this document for use as a technical specification is expressly prohibited.

The notice is:

"Copyright (c) [\$date-of-document] Eclipse Foundation. This software or document includes material copied from or derived from [title and URI of the Eclipse Foundation specification document]."

Disclaimers

THIS DOCUMENT IS PROVIDED "AS IS," AND THE COPYRIGHT HOLDERS AND THE ECLIPSE FOUNDATION MAKE NO REPRESENTATIONS OR WARRANTIES, EXPRESS OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NON-INFRINGEMENT, OR TITLE; THAT THE CONTENTS OF THE DOCUMENT ARE SUITABLE FOR ANY PURPOSE; NOR THAT THE IMPLEMENTATION OF SUCH CONTENTS WILL NOT INFRINGE ANY THIRD PARTY PATENTS, COPYRIGHTS, TRADEMARKS OR OTHER RIGHTS.

THE COPYRIGHT HOLDERS AND THE ECLIPSE FOUNDATION WILL NOT BE LIABLE FOR ANY DIRECT, INDIRECT, SPECIAL OR CONSEQUENTIAL DAMAGES ARISING OUT OF ANY USE OF THE

DOCUMENT OR THE PERFORMANCE OR IMPLEMENTATION OF THE CONTENTS THEREOF.

The name and trademarks of the copyright holders or the Eclipse Foundation may NOT be used in advertising or publicity pertaining to this document or its contents without specific, written prior permission. Title to copyright in this document will at all times remain with copyright holders.

Jakarta Data

Chapter 1. Introduction

The Jakarta Data specification provides an API for easier data access. A Java developer can split the persistence from the model with several features, such as the ability to compose custom query methods on a Repository interface where the framework will implement it.

There is no doubt about the importance of data around the application. We often talk about a stateless application, where we delegate the application's state to the database.

Dealing with a database is one of the biggest challenges within a software architecture. In addition to choosing one of several options on the market, it is necessary to consider the persistence integrations. Jakarta Data makes life easier for Java developers.

1.1. Goals

Jakarta Data works in a tight integration between Java and a persistence layer, where it has the following specification goals:

- Be a persistence agnostic API. Therefore, through abstractions, it will connect different types of databases and storage sources.
- Be a pluggable and extensible API. Even when the API won't support a particular behavior of a storage engine, it might provide an extensible API to make it possible.

1.2. Non-Goals

As with any software component, these decisions come with trade-offs and the following non-goals:

- Provide specific features of Jakarta Persistence, Jakarta NoSQL, etc. Those APIs have their own specifications.
- Replace the Jakarta Persistence or Jakarta NoSQL specifications. Indeed, Jakarta Data will work as a complement to these specifications as an agnostic API.

1.3. Conventions

1.4. Jakarta Data Project Team

This specification is being developed as part of Jakarta Data project under the Jakarta EE Specification Process. It is the result of the collaborative work of the project committers and various contributors.

1.4.1. Project Leads

- [Nathan Rauh](#)
- [Otavio Santana](#)

1.4.2. Committers

- [Denis Stepanov](#)
- [Dmitry Kornilov](#)
- [Emily Jiang](#)
- [Graeme Rocher](#)
- [James Krueger](#)
- [James Stephens](#)
- [Michael Redlich](#)
- [Nathan Rauh](#)
- [Otavio Santana](#)
- [Werner Keil](#)

1.4.3. Mentor


- [Dmitry Kornilov](#)

1.4.4. Contributors

The complete list of Jakarta Data contributors may be found [here](#).

Chapter 2. Repository

In Domain-Driven Design (DDD) the repository pattern encapsulates the logic required to access data sources. The repository pattern consolidates data access functionality, providing better maintainability and decoupling the infrastructure or technology used to access databases from the domain model layer.



This pattern focuses on the closest proximity of entities and hides where the data comes from.

The Repository pattern is a well-documented way of working with a data source. In the book *Patterns of Enterprise Application Architecture*, Martin Fowler describes a repository as follows:

A repository performs the tasks of an intermediary between the domain model layers and data mapping, acting in a similar way to a set of domain objects in memory. Client objects declaratively build queries and send them to the repositories for answers. Conceptually, a repository encapsulates a set of objects stored in the database and operations that can be performed on them, providing a way that is closer to the persistence layer. Repositories also support the purpose of separating, clearly and in one direction, the dependency between the work domain and the data allocation or mapping.

It also becomes very famous in *Domain-Driven Design: Tackling Complexity in the Heart of Software* by Eric Evans.

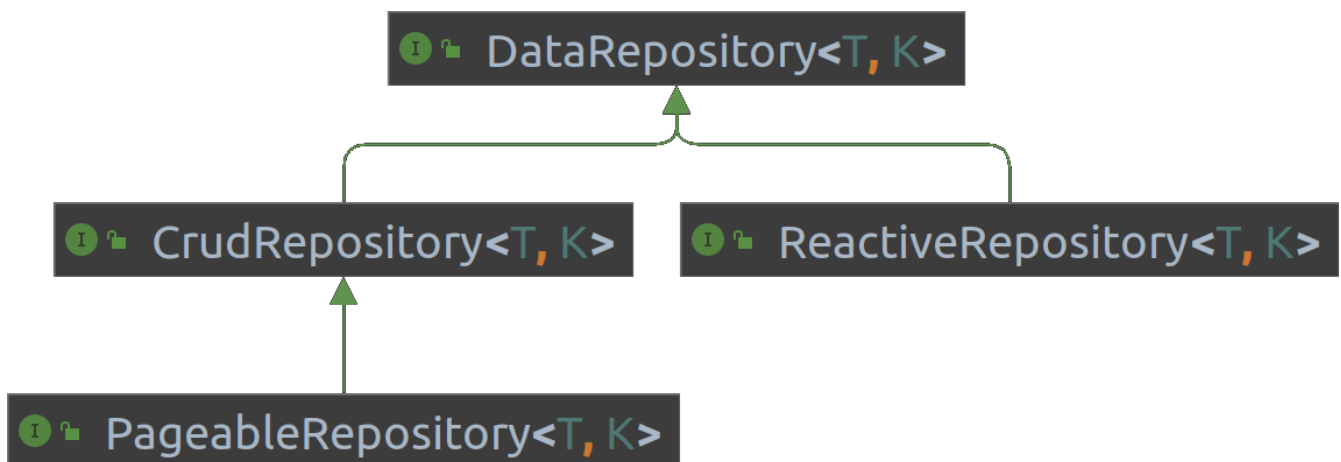
2.1. Repositories on Jakarta Data

A repository abstraction aims to significantly reduce the boilerplate code required to implement

data access layers for various persistence stores.

The parent interface in Jakarta Data repository abstraction is `DataRepository`.

By default, Jakarta Data has support for three interfaces. However, the core is extensible. Therefore, a provider might extend one or more interfaces to a specific data target.



- Interface to generic CRUD operations on a repository for a specific type. This one we can see more often on several Java implementations.
- Interface with generic CRUD operations using the pagination feature.
- Interface for generic CRUD operations on a repository for a specific type. This repository follows reactive paradigms.

From the Java developer perspective, create an interface that is annotated with the `Repository` annotation and optionally extends one of the built-in repository interfaces.

So, given a `Product` entity where the ID is a `long` type, the repository would be:

```
@Repository
public interface ProductRepository extends CrudRepository<Product, Long> {

}
```

There is no nomenclature restriction to make mandatory the `Repository` suffix. Such as, you might represent the repository of the Car's entity as a `Garage` instead of `CarRepository`.

```
@Repository
public interface Garage extends CrudRepository<Car, String> {

}
```

2.2. Entity Classes

Entity classes are simple Java objects with fields or accessor methods designating each entity property.

TODO Write the remainder of this section after we determine if a common entity model can be created, or if Jakarta Data will define its own entity model, or if we will reuse the entity models from Jakarta Persistence and Jakarta NoSQL.

2.2.1. Entity Property Names

Within an entity, property names must be unique ignoring case and are either computed from the field or accessor method name or by other means defined by the entity model, such as through the `@Id` or `@Column` annotations, which take precedence if present on the field or accessor method. The entity property that is chosen as the id is referred to in repository methods by the name `Id` or optionally by its other designated name, for example if one is assigned by an `@Id` or `@Column` annotation, with the latter being used to refer to the id property in query language. Entity properties which are not the id give precedence to the `@Column` name if specified or otherwise the field or accessor method name in both name-pattern based repository query methods and query language. Entity property names that are used in composing name-pattern based repository query methods must not contain reserved words.

2.3. Queries Methods

In Jakarta Data, besides finding by an ID, custom queries can be written in two ways:

- Through Query annotation: It will create a method annotated with the `@Query` with the query.
- Using method by query convention: Using some pattern vocabulary will provide a query.



Due to the variety of data sources, those resources might not work; it varies based on the Jakarta Data implementation and the database engine, which can provide queries on more than a Key or ID or not, such as a Key-value database.

2.3.1. Using the Query Annotation

The Query's annotation will support a search expression as a String. The specification won't define the query syntax, which might vary between vendors and data sources, such as SQL, JPA-QL, Cypher, CQL, etc.

```
@Repository
public interface ProductRepository extends CrudRepository<Product, Long> {
    @Query("SELECT p FROM Products p WHERE p.name=?1") // example in JPQL
    Optional<Product> findByName(String name);
}
```

Jakarta Data also includes the `Param` annotation to define a binder annotation, where as with the query expression, each vendor will express the syntax freely such as `?`, `@`, etc..

```
@Repository
public interface ProductRepository extends CrudRepository<Product, Long> {
    @Query("SELECT p FROM Products p WHERE p.name=:name") // example in JPQL
}
```

```
Optional<Product> findByName(@Param("name") String name);  
}
```

2.3.2. Query by Method

The Query by method mechanism allows for creating query commands by conventions.

E.g.:

```
@Repository  
public interface ProductRepository extends CrudRepository<Product, Long> {  
  
    List<Product> findByName(String name);  
  
    @OrderBy("price")  
    List<Product> findByNameLike(String namePattern);  
  
    @OrderBy(value = "price", descending = true)  
    List<Product> findByNameLikeAndPriceLessThan(String namePattern, float priceBelow);  
  
}
```

The parsing query method name has two parts: the subject and the property.

The first part defines the query's subject or condition, and the second the condition value; both forms the predicate.

A predicate can refer only to a direct property of the managed entity. We also have the option to handle entities with another class on them.

Assume an Order entity has an Address with a Zipcode. In that case, the access is `order.address.zipCode`.

The resolution algorithm starts by interpreting the whole part (AddressZipCode) as the property and checks the domain class for a property with that name (uncapitalized). If the algorithm succeeds, it uses that property.

```
@Repository  
public interface OrderRepository extends CrudRepository<Order, Long> {  
  
    Stream<Order> findByAddressZipCode(ZipCode zipCode);  
  
}
```

Although this should work for most cases, to resolve this ambiguity, you can use `_` inside your method name to manually define traversal points.

```
@Repository
public interface OrderRepository extends CrudRepository<Order, Long> {

    Stream<Order> findByAddress_ZipCode(ZipCode zipCode);

}
```



Define as a priority following standard Java naming conventions, camel case, using underscore as the last instance resource.

2.3.2.1. Query Methods Keywords

The following table lists the subject keywords generally supported by Jakarta Data.

Keyword	Description
findBy	General query method returning the repository type.
deleteBy	Delete query method returning either no result (void) or the delete count.
countBy	Count projection returning a numeric result.
existsBy	Exists projection, returning typically a boolean result.

Jakarta Data implementations often support the following list of predicate keywords. It might not work because some keywords listed here might not be supported in a particular store.

Keyword	Description	Method signature Sample
And	The and operator.	findByNameAndYear
Or	The or operator.	findByNameOrYear
Between	Find results where the property is between the given values	findByDateBetween
Empty	Find results where the property is an empty collection or has a null value.	deleteByPendingTasksEmpty
LessThan	Find results where the property is less than the given value	findByAgeLessThan
GreaterThan	Find results where the property is greater than the given value	findByAgeGreaterThan
LessThanEqual	Find results where the property is less than or equal to the given value	findByAgeLessThanEqual

Keyword	Description	Method signature Sample
GreaterThanOrEqualTo	Find results where the property is greater than or equal to the given value	findByAgeGreaterThanOrEqualTo
Like	Finds string values "like" the given expression	findByTitleLike
IgnoreCase	Requests that string values be compared independent of case for query conditions and ordering.	findByStreetNameIgnoreCaseLike
In	Find results where the property is one of the values that are contained within the given list	findByIdIn
Null	Finds results where the property has a null value.	findByYearRetiredNull
True	Finds results where the property has a boolean value of true.	findBySalariedTrue
False	Finds results where the property has a boolean value of false.	findByCompletedFalse
OrderBy	Specify a static sorting order followed by the property path and direction of ascending.	findByNameOrderByAge
OrderBy____Desc	Specify a static sorting order followed by the property path and direction of descending.	findByNameOrderByAgeDesc
OrderBy____Asc	Specify a static sorting order followed by the property path and direction of ascending.	findByNameOrderByAgeAsc
OrderBy____(Asc Desc)*(Asc Desc)	Specify several static sorting orders	findByNameOrderByAgeAscNameDescYearAsc

2.4. Special Parameter Handling

Jakarta Data also supports particular parameters to define pagination and sorting.

Jakarta Data recognizes, when specified on a repository method after the query parameters, specific types, like **Limit**, **Pageable**, and **Sort**, to dynamically apply limits, pagination, and sorting to queries. The following example demonstrates these features:

```
@Repository
public interface ProductRepository extends CrudRepository<Product, Long> {
```

```
List<Product> findByName(String name, Pageable pageable);

List<Product> findByNameLike(String pattern, Limit max, Sort... sorts);

}
```

You can define simple sorting expressions by using property names.

```
Sort name = Sort.asc("name");
```

You can combine sorting with a starting page and maximum page size by using property names.

```
Pageable pageable = Pageable.ofSize(20).page(1).sortBy(Sort.desc("price"));
first20 = products.findByNameLike(name, pageable);
```

2.5. Keyset Pagination

Keyset pagination aims to reduce missed and duplicate results across pages by querying relative to the observed values of entity properties that constitute the sorting criteria. Keyset pagination can also offer an improvement in performance because it avoids fetching and ordering results from prior pages by causing those results to be non-matching. A Jakarta Data provider appends additional conditions to the query and tracks keyset values automatically when `KeysetAwareSlice` or `KeysetAwarePage` are used as the repository method return type. The application invokes `nextPageable` or `previousPageable` on the keyset aware slice or page to obtain a `Pageable` which keeps track of the keyset values.

For example,

```
@Repository
public interface CustomerRepository extends CrudRepository<Customer, Long> {
    KeysetAwareSlice<Customer> findByZipcodeOrderByLastNameAscFirstNameAscIdAsc(
        int zipcode, Pageable pageable);
}
```

You can obtain the next page with,

```
for (Pageable p = Pageable.ofSize(50); p != null; ) {
    page = customers.findByZipcodeOrderByLastNameAscFirstNameAscIdAsc(55901, p);
    ...
    p = page.nextPageable();
}
```

Or you can obtain the next (or previous) page relative to a known entity,

```
Customer c = ...
Pageable p = Pageable.ofSize(50).afterKeyset(c.lastName, c.firstName, c.id);
page = customers.findByZipcodeOrderByLastNameAscFirstNameAscIdAsc(55902, p);
```

The sort criteria for a repository method that performs keyset pagination must uniquely identify each entity and must be provided by one of:

- **OrderBy** name pattern of the repository method (as in the examples above).
- **@OrderBy** annotation(s) on the repository method.
- **Sort** parameters of the **Pageable** that is supplied to the repository method.

If either of the first are specified, the **Sort** list of the **Pageable** must be empty or unspecified.

2.5.1. Example of Appending to Queries for Keyset Pagination

Without keyset pagination, a Jakarta Data provider that is based on Jakarta Persistence might compose the following JPQL for the **findByZipcodeOrderByLastNameAscFirstNameAscIdAsc** repository method from the prior example:

```
SELECT o FROM Customer o WHERE (o.zipCode = ?1)
                                ORDER BY o.lastName ASC, o.firstName ASC, o.id ASC
```

When keyset pagination is used, the keyset values from the **Cursor** of the **Pageable** are available as query parameters, allowing the Jakarta Data provider to append additional query conditions. For example,

```
SELECT o FROM Customer o WHERE (o.zipCode = ?1)
                                AND ( (o.lastName > ?2)
                                      OR (o.lastName = ?2 AND o.firstName > ?3)
                                      OR (o.lastName = ?2 AND o.firstName = ?3 AND o.id >
?4)
                                )
                                ORDER BY o.lastName ASC, o.firstName ASC, o.id ASC
```

2.5.2. Avoiding Missed and Duplicate Results

Because searching for the next page of results is relative to a last known position, it is possible with keyset pagination to allow some types of updates to data while pages are being traversed without causing missed results or duplicates to appear. If you add entities to a prior position in the traversal of pages, the shift forward of numerical position of existing entities will not cause duplicates entities to appear in your continued traversal of subsequent pages because keyset pagination does not query based on a numerical position. If you remove entities from a prior position in the traversal of pages, the shift backward of numerical position of existing entities will not cause missed entities in your continued traversal of subsequent pages because keyset pagination does not query based on a numerical position.

Other types of updates to data, however, will cause duplicate or missed results. If you modify entity properties which are used as the sort criteria, keyset pagination cannot prevent the same entity from appearing again or never appearing due to the altered values. If you add an entity that you previously removed, whether with different values or the same values, keyset pagination cannot prevent the entity from being missed or possibly appearing a second time due to its changed values.

2.5.3. Restrictions on use of Keyset Pagination

- The repository method signature must return `KeysetAwareSlice` or `KeysetAwarePage`. A repository method with return type of `KeysetAwareSlice` or `KeysetAwarePage` must raise `UnsupportedOperationException` if the database is incapable of keyset pagination.
- The repository method signature must accept a `Pageable` parameter.
- Sort criteria must be provided and should be minimal.
- The combination of provided sort criteria must uniquely identify each entity.
- Page numbers for keyset pagination are estimated relative to prior page requests or the observed absence of further results and are not accurate. Page numbers must not be relied upon when using keyset pagination.
- Page totals and result totals are not accurate for keyset pagination and must not be relied upon.
- A next or previous page can end up being empty. You cannot obtain a next or previous `Pageable` from an empty page because there are no keyset values relative to which to query.
- A repository method that is annotated with `@Query` and performs keyset pagination must omit the `ORDER BY` clause from the provided query and instead must supply the sort criteria via `@OrderBy` annotations or `Sort` parameters of `Pageable`. The provided query must end with a `WHERE` clause to which additional conditions can be appended by the Jakarta Data provider. The Jakarta Data provider is not expected to parse query text that is provided by the application.

2.5.4. Keyset Pagination Example with Sorts

Here is an example where an application uses `@Query` to provide a partial query to which the Jakarta Data provider can generate and append additional query conditions and an `ORDER BY` clause.

```
@Repository
public interface CustomerRepository extends CrudRepository<Customer, Long> {
    @Query("SELECT o FROM Customer o WHERE (o.totalSpent / o.totalPurchases > ?1)")
    KeysetAwareSlice<Customer> withAveragePurchaseAbove(float minimum, Pageable
    pagination);
}
```

Example traversal of pages:

```
for (Pageable p = Pageable.ofSize(25).sortBy(Sort.desc("yearBorn"), Sort.asc("name"),
Sort.asc("id"));
    p != null; ) {
    page = customers.withAveragePurchaseAbove(50.0f, p);
    ...
}
```

```
p = page.nextPageable();  
}
```